

TxBench-PP: Analyzing AI Agent Performance on Preclinical Pharmacology

Hannah Le^{1,*}, Ramesh Ramasamy^{1,*}, Alex Urrutia^{1,*}, Mahsa Yazdani^{1,*}, Tim Proctor¹, Kenny Workman¹

¹LatchBio, San Francisco, CA, USA

*These authors contributed equally.

Correspondence: kenny@latch.bio

ABSTRACT


AI agents promise to accelerate drug discovery by compressing interpretation and decision-making loops, but practical deployment requires trusted evaluation on realistic program decisions. We introduce TherapeuticsBench Preclinical Pharmacology (TxBench-PP), a verifiable benchmark for small-molecule preclinical pharmacology and the first focused slice of a broader TherapeuticsBench effort across drug-discovery stages and therapeutic modalities. TxBench-PP tests whether agents can recover accurate conclusions from real world assay data rather than memorized facts from literature. The benchmark contains 100 evaluations indexed by program stage, assay evidence, and task type, spanning mechanism-of-action and pharmacodynamic reasoning, compound-target engagement, causal target validation, developability and safety, and translational efficacy. Agents receive realistic workflow snapshots, inspect files in a coding environment, and return structured answers graded deterministically. Across 16 model-harness configurations, comprising 11 models and 4,800 trajectories, no system reliably recovered preclinical pharmacology decisions. The strongest configuration, Claude Opus 4.8 / Pi, passed 59.3% of trajectories, followed by GPT-5.5 / Pi at 55.3%.

TxBench-PP: Overview

TxBench-PP is the small-molecule preclinical pharmacology slice of TherapeuticsBench. It does not try to cover the full drug-discovery lifecycle. Instead, it focuses on local decisions that determine whether a target or compound should advance, hold, or fail: whether the intended biology is engaged, whether the assay evidence supports a mechanism, whether exposure can cover the relevant potency range, and whether a safety signal blocks progression.

TherapeuticsBench is the broader roadmap across development phases and therapeutic modalities. The schematic below places TxBench-PP within that map. The stage labels identify the main evidence type tested by each evaluation; they are not a required chronology for every program.


A TxBench-PP covers selected small-molecule preclinical decisions



Current release: small-molecule preclinical pharmacology task groups. Future: other modalities and full program stories.

B Endpoint grading

100 tasks x 3 attempts per model-harness pair



C Cost-performance frontier




Figure 1: TxBench-PP scope and topline performance. A, TxBench-PP is a focused small-molecule preclinical pharmacology slice of a broader TherapeuticsBench roadmap. The dotted boundary marks the task groups covered in this release; biologics, oligonucleotides, and full program histories are out of scope. B, Endpoint pass rates across 16 model-harness configurations, with 95% confidence intervals computed over task-level pass rates. C, Mean cost per task against endpoint pass rate.

Introduction

While experiments are rate-limited by natural processes, human decisions and organizational consensus often make up significant components of program timelines in drug discovery. Agents in drug discovery promise to accelerate discovery, development, and translation by compressing these interpretation and decision-making loops. Specific knowledge from constituent steps of the drug-development pipeline can be encoded into frontier models through post-training, tool use, and domain-specific scaffolding, then deployed into practical data interpretation workflows.

However, the practical use of agentic systems in industrial workflows requires standard and trusted methods of evaluating performance. This is especially challenging because the drug-discovery ecosystem is a sprawling landscape of assay categories, development stages, therapeutic modalities, and decision types. Benchmarks must therefore measure realistic tasks while providing focused treatment of the many local scientific judgments that tile the biotech ecosystem.

Preclinical small-molecule pharmacology is one important piece of this landscape. Potency, target engagement, mechanism-of-action evidence, pharmacodynamic response, exposure, safety, and in vivo efficacy are converted into practical decisions about whether to advance, hold, repeat, or kill a candidate. These decisions sit upstream of costly medicinal chemistry, animal studies, and human experiments. They are easy to get wrong and often depend on fragmented expert judgment.

We introduce TxBench-PP, a verifiable benchmark for evaluating AI agents on small-molecule preclinical pharmacology. TxBench-PP is best understood as a focused benchmark within the broader TherapeuticsBench roadmap rather than a complete evaluation of drug discovery holistically. Each evaluation supplies realistic workflow artifacts from a local preclinical decision point and asks the agent to return a structured, deterministically gradable answer. Importantly, the benchmark tests if agents can recover decisions from supplied data instead of leaning on well-understood mechanisms and literature knowledge, with many tasks intentionally designed to penalize systems that do this.

TxBench-PP contains 100 evaluations indexed by program stage, assay type, and task structure. Tasks span mechanism-of-action and pharmacodynamic reasoning, compound-target engagement, causal target validation, developability and safety, and translational efficacy. Other parts of the therapeutic ecosystem, including clinical program stages and non-small-molecule modalities, are left to future TherapeuticsBench efforts. The benchmark offers a measuring stick by which scientists can evaluate agentic systems on local preclinical pharmacology decisions toward practical deployment in drug-discovery programs.

Benchmark Construction

Agents are evaluated on realistic preclinical pharmacology decisions over 100 evaluations spanning eight program stages: disease and model context, screening and hit confirmation, drug response, causal target validation, compound-target characterization, MoA/PD, developability/safety/PK, and translational efficacy. Human genetics tasks are reserved for future benchmarking work. Across the benchmark, evaluations draw on assay types including drug-response screens, genetic perturbation, transcriptomic and proteomic profiling, chemoproteomic and live-cell target engagement, imaging, DMPK/ADMET panels, toxicogenomics, and in vivo efficacy or safety studies.

The benchmark is built around recurring local decisions in small-molecule programs. Candidate tasks begin as specific analysis points where assay data might be turned into decisions. Examples include identifying a pathway mechanism from dose-resolved phosphoproteomics, separating true chemoproteomic target engagement from complex contaminants, and interpreting in vivo safety data.

Each evaluation includes the data artifacts needed for that decision, metadata needed to interpret the files, a task description, and a structured answer. Task context is calibrated to approximate what a scientist would know at the point of analysis. We avoid prescribing a method unless the method is part of the scientific decision. The intended answer is a result supported by the supplied data, not some fact or mechanism well understood in the literature. Many tasks are designed to intentionally trap systems that overfit on trained knowledge without empirical exploration.

Candidate evaluations are refined through manual review and model trajectory inspection. We remove tasks when the answer can be recovered by an obvious shortcut, when the prompt over-specifies the analysis, or when plausible analysis choices produce answers that break grading assumptions. The retained tasks are tagged by program stage, assay type, and task structure.

Grading uses deterministic functions over structured final answers. Depending on the task, graders check label sets, numerical tolerances, rank-order relationships, categorical choices, or all required fields. Manual trajectory inspection is part of benchmark construction. We track decision-blocking gaps that should change the program decision if recognized: tox-species mismatch, insufficient free-drug coverage, hook effects, contaminated chemoproteomic signals, survivor bias in pathology sampling, and related failure points. These annotations are not exposed to agents; they are used to interpret failures after endpoint grading and to guide future benchmark updates.

Benchmark Anatomy

TxBench-PP is organized around the anatomy of a practical preclinical pharmacology decision. Each evaluation is tagged along three axes: where the decision sits in a small-molecule program, the assay type that generated the data, and the structure of the task.


The program-stage axis follows a rough temporal ordering: disease and model context, human genetic target support, causal target validation, screening and hit confirmation, compound-target characterization, MoA/PD, drug response, developability/safety/PK, and translational efficacy. Human genetic target support is reserved for future benchmarking work because adequately measuring this category requires a separate focused

benchmark. The order gives readers context for task progression; it does not imply that every small-molecule program proceeds through this strict sequence in practice.

The assay and task axes describe the actual work required. Assay labels collapse source kits into measurement families such as drug-response screens, protein or PTM measurements, genetic perturbation, target engagement, molecular state, safety profiling, and in vivo experiments. Task labels describe broad families of decisions: mechanism inference, QC, dependency calls, differential response, safety assessment, imaging, hit prioritization, and related judgments.

Benchmark anatomy

Program stage locates the decision; assay type and task structure describe the data and work required.



* reserved for future benchmarking work

Task counts show categories with ≥ 4 evaluations plus the smaller-category tail.

Figure 2: Benchmark anatomy. A, Program stage labels place evaluations along a roughly ordered small-molecule evidence spine; the asterisk marks a stage reserved for future benchmarking work. B, Assay-type labels summarize the measurement families represented in the source data. C, Task-structure labels describe the operation or judgment required for a passing answer.

Results

No agent reliably recovers preclinical pharmacology decisions

We evaluated 16 model-harness configurations, comprising 11 models across three agent harnesses, on 100 preclinical pharmacology tasks. Each configuration was run three independent times per task, yielding 4,800 agent trajectories. A trajectory was scored as passing only if its submitted answer passed every task-specific grader. We report pass rate as the fraction of passing trajectories, with 95% confidence intervals computed across tasks.

Performance at the top was clustered rather than clearly separated. The strongest configuration, Claude Opus 4.8

on Pi, achieved a pass rate of 59.3% (178/300 trajectories; 95% CI, 51.0–67.6), followed by GPT-5.5 on Pi at 55.3% (166/300; 47.0–63.6), Claude Opus 4.8 on Claude Code at 54.7% (164/300; 45.9–63.4), and Gemini-3.5-Flash on Pi at 51.3% (154/300; 42.9–59.8). These performances overlap substantially, indicating a broad leading tier rather than a decisive winner. More importantly, even the best-performing configuration remained below 60%, leaving more than 40% of task attempts unsolved.

Thus, across providers, model generations, and agent harnesses, no system approached the reliability expected for preclinical decision-making. The absence of a clear leader, together with a performance ceiling near 60%, suggests that current agent systems still do not reliably recover preclinical pharmacology decisions from the provided data.




Figure 3: Pass rates and reliability across model-harness configurations. (A) Strict pass rate for each of the 16 model-harness configurations across all 4,800 agent trajectories, with 95% bootstrap confidence intervals computed over the 100 tasks. Each trajectory was scored pass only when its submitted answer satisfied every task-specific grader. Alternate harnesses (Claude Code, Codex) are shown as offset points for the models that ran them. (B) Task-level reliability for the Pi harness, partitioning the 100 evaluations for each model by how many of the three independent rounds passed.

Model / harness	Endpoint pass rate	$\geq 1/3$	$\geq 2/3$	3/3
Claude Opus 4.8 / Pi	59.3% (178/300; 51.1–67.6)	72/100	65/100	41/100
GPT-5.5 / Pi	55.3% (166/300; 47.0–63.6)	74/100	51/100	41/100
Claude Opus 4.8 / Claude Code	54.7% (164/300; 45.9–63.4)	65/100	59/100	40/100
Gemini 3.5 Flash / Pi	51.3% (154/300; 42.9–59.8)	67/100	52/100	35/100
GPT-5.4 / Pi	49.7% (149/300; 41.3–58.0)	67/100	48/100	34/100
Claude Opus 4.7 / Pi	49.3% (148/300; 40.4–58.3)	60/100	50/100	38/100
GPT-5.5 / OpenAI Codex	47.3% (142/300; 39.6–55.1)	70/100	46/100	26/100
GPT-5.4 / OpenAI Codex	46.7% (140/300; 38.8–54.5)	68/100	46/100	26/100
Claude Opus 4.6 / Pi	44.7% (134/300; 36.1–53.2)	59/100	45/100	30/100
Claude Opus 4.7 / Claude Code	44.0% (132/300; 35.9–52.1)	62/100	45/100	25/100
Claude Opus 4.6 / Claude Code	41.3% (124/300; 32.9–49.7)	55/100	43/100	26/100
Gemini 3.1 Pro / Pi	40.0% (120/300; 31.5–48.5)	55/100	37/100	28/100
Claude Sonnet 4.6 / Pi	36.0% (108/300; 28.4–43.6)	56/100	35/100	17/100
Kimi K2P6 / Pi	29.7% (89/300; 22.3–37.0)	48/100	26/100	15/100
Grok 4.20 reasoning / Pi	19.7% (59/300; 13.4–25.9)	34/100	18/100	7/100
Grok 4.3 / Pi	18.3% (55/300; 11.9–24.8)	30/100	15/100	10/100

Table 1: Pass rates by model–harness configuration. Pass rates are reported as percentages, with the numerator and denominator shown as passing trajectories out of 300 total trajectories per configuration, together with 95% confidence intervals. Confidence intervals were computed across task-level mean pass rates over the 100 preclinical pharmacology tasks. The final three columns report the number of tasks, out of 100, for which at least one, at least two, or all three independent attempts passed.

Most failures map to a compact failure-mode taxonomy

To understand why agents fail, we manually reviewed every failing (model, evaluation) run on the Pi harness and labeled each against the evaluation’s documented scientific intent, yielding 1,834 labeled failures. We classify each failure along two complementary axes: the *scientific error*, which captures what is wrong with the answer, and the *behavioral cause*, which captures how the model arrived there (Figure 4).

The scientific axis comprises five recurring error families (Figure 4A, restricted to the $n = 1,436$ failures that contain a scientific error). The largest, flawed method, resolves into two main sub-modes and a small tail. A *skipped-QC* error (17%) omits a required preprocessing or quality-control step: in one case, a model aggregated CRISPR guides to gene-level scores without first dropping guides that map to multiple loci, so paralog-family artifacts flooded the essential-gene list. A *wrong-reference-frame* error (21%) runs the pipeline correctly but compares against the wrong baseline, comparator, or normalization. The small *wrong-assay-layer* tail (2%) reads the wrong readout from the right experiment. On a proteasome-kinetics task asking which compound degrades substrate faster, the answer lives in the ubiquitin-remnant layer of a dual-readout mass-spectrometry run. A model that answers from the phosphoproteome layer of the same run reaches a different conclusion on what is a different question.

Beyond flawed method, four families remain. A *calibration* error (31%) draws a cutoff a domain expert would set differently. For instance, in one competition-binding panel, models carried forward all 33 nonzero kinases when only the 8 strong binders above a clear affinity gap were supported. A *perception* error (10%) misreads a raw image, such as counting apoptotic debris as viable organoids or mis-segmenting two-channel colocalization. *Prior over data* (10%) answers from textbook or famous biology against the staged evidence: one model computed the

strongest data-supported biomarker—a vanadium compound correlated with SLC26A2 at $r = -0.48$ —then discarded it for the textbook elesclomol-FDX1 cuproptosis pair. *Over-reading* (10%) states a stronger claim than the data supports, such as reporting tumor regression from a sub-baseline median when the per-animal measurements show only growth arrest. Together, method and calibration account for 71% of scientific errors.

The behavioral axis records how each failure happened, across all $n = 1,834$ runs (Figure 4B). Two-thirds (67%) are *engaged, analytically wrong*: the model did the work in good faith and made an analytical error. The remaining third splits into one evidence-handling failure and three process failures. *Prior-driven override* (6%), the lone evidence-handling case, engages the data and often computes the supporting result, then defers to a textbook or famous-biology prior. *Weak commitment* (9%) performs the right analysis but submits a final answer that does not match it, a mismatch that appears even in the strongest models and points to a missing final-answer check. *Abandoned-correct* (12%) has a passing result within reach and then commits to a worse, more canonical one such as a published figure or standard global method. Like prior-driven override, it occurs most often in the longest runs. *Under-investment* (6%) reaches an answer with little tool use, submitting a plausible call before running the per-dose computation or QC the task requires.

The two axes distribute differently across models (Figure 4C,D). The behavioral mix varies by model. Under-investment is concentrated in the two weakest models, at 32% of grok-4.3’s failures and 17% of grok-4.20-reasoning’s against 0–1% for the others. Abandoned-correct is spread across the frontier at 7–17%, highest in gpt-5.5 and gemini-3.1-pro (17% each). The scientific-error mix is nearly constant, with method and calibration accounting for 58–77% of every model’s scientific errors regardless of pass rate.




Figure 4: A two-axis failure taxonomy. All 1,834 failing (model, evaluation) runs on the Pi harness, labeled against each evaluation’s scientific intent and classified along two complementary axes. **(A)** The scientific-error axis (what is wrong with the answer) is dominated by method and calibration errors (71% combined), which are general analytical skills rather than domain knowledge; the method family resolves into a skipped-QC/preprocessing sub-mode and a larger wrong-reference-frame sub-mode. **(B)** The behavioral-cause axis (how the model failed) shows that a third of failures are non-analytical: a prior-driven override of the model’s own evidence, plus three process breakdowns (weak commitment, abandoned-correct, under-investment). **(C)** The behavioral mix is a model signature: under-investment dominates only the weakest model, while the strongest models fail almost entirely through engaged, analytically wrong reasoning. **(D)** The scientific-error mix, by contrast, is largely model-invariant: method and calibration errors dominate every model regardless of overall accuracy, so the *kinds* of scientific mistakes are shared across the frontier even as their rate falls (each bar normalized within that model’s scientific-error failures). A scientific-error × behavioral-cause cross-tabulation is provided in the supplement (fig_fm_crosstab).

The dominant failure mode shifts characteristically across pipeline stages

Pass rates range from 27% at the hardest stage to 55% at the easiest stage (Figure 5). Walking the arc in order shows that each stage has characteristic strengths and weaknesses. We look at a representative easy and hard evaluation for each stage, with the fraction of passing runs in parentheses.

Overall, the ordering of models at each stage closely mirrors their overall performance (Figure 5C; mean Spearman $\rho = 0.80$ across stages, ranging from 0.56 at S7 to 0.92 at S8), with the two Grok models ranking lowest at every stage. Three features stand out against this baseline. The two highest cells in the matrix are gemini-3.5-flash and gpt-5.5 at drug response, at 87% (95% CI 70–95, $n = 30$) and 86% (69–95, $n = 29$), well above their own overall rates of 52% and 55%. The clearest stage effect is within a single model. gemini-3.5-flash falls from 87% at S3 to 28% (16–44, $n = 36$) at causal target validation. The widest within-family spread is at hit screening, where claude-opus-4-8 scores 50% (31–69, $n = 24$) and the older claude-opus-4-6 12% (4–31, $n = 24$).

Primary screening and hit prioritization (27%) is the hardest stage. Models are competent at the mechanical steps, such as robustly collapsing noisy replicates (PRISM05, 58%). They fail by choosing the wrong reference frame for QC. On outlier-pool QC (PRISM01, 18%) they flag whole well-by-pool groups by an aggregate median residual and exclude most of the cytotoxic positive-control wells, discarding real biology as a technical artifact. On an oncology drugs repurposing prioritization task (PRISM_S2_11, 0%) they rank by breadth of killing and nominate a pan-cytotoxic compound rather than a selective hit.

Drug response (55%) is one of the easier stages but exposes a consistent weakness in unsupervised mechanism recovery. Models do well on direct evidence checks, such as separating a causal mediator from a correlated biomarker (PRISM_S4_06, 76%). They fail when asked to cluster compounds by mechanism from their response profiles. Across the embedding tasks they use a potency-dominated representation, a linear PCA or a Euclidean metric on raw profiles, so the leading component encodes potency and the intended drug class never separates from the rest. The same error recurs across statin collapse (PRISM_UMAP_01, 12%), oncology MoA recovery (PRISM_UMAP_04, 48%), and glucocorticoid substructure (PRISM_UMAP_03, 69%), so it is a property of the stage rather than a single task.

Causal target validation (32%) is the second-hardest stage and the one that most rewards data hygiene and statistical rigor. Models recover genuine dependencies from a clean screen (HIT_05 replicate discordance, 69%). They routinely skip the library hygiene that CRISPR screens require. On multi-target guide contamination (GG_01, 16%) they aggregate guides to gene level without dropping multi-locus sgRNAs, so paralog-family false positives flood the dependency list. On the selective-set call (SEL_01, 6%) they ship thousands of

genes behind a count-only filter with no effect-size floor.

Mechanism of action and PD (37%) is the largest stage, with 21 evaluations. Models reliably recover known program structure, such as assigning cells to the correct cell-cycle phase (clust_04, 84%). Their recurring failure is over-reaching on the mechanism, treating marginal or partial evidence as a strong, broad effect. On the antiinfective task (ANTIINF_01, 0%) they read near-noise pathway scores ($p \approx 0.05$) as genuinely engaged mechanisms. On the mTOR fingerprint (TXBX_DM_v15E, 6%) they correctly call mTOR the primary target but then count phosphosites from too broad a substrate panel, sweeping in adjacent AKT-axis sites instead of restricting to the canonical mTOR set, which inflates the affected-site count to almost twice what the data supports (59 versus 31).

Compound-target engagement (43%) asks whether a compound truly hits its target and how selectively, and the bottleneck is calibration, not detection (67% of the stage's scientific errors). Models detect the binding signal reliably, for example confirming a clean direct target by thermal shift (thermal_shift_shared_direct_target, 91%). What they lack is a principled cutoff for how much engagement counts, so they err in both directions. On a kinase-binding panel (TXBX_KB_v11_DA1, 0%) they accept every weak binder as a real hit, and on a control task (CTRL01, 21%) they invent a stringency cutoff the data never justified and discard five genuinely engaged targets.

Developability and safety (43%) depends on judging a signal against the right reference, and that is where models slip. They make a confident call when the standard is unambiguous, such as the hepatotoxicity verdict (TXBX_TG_v11_A, 91%). On the hERG safety-margin task (de_14) they size the margin against total drug concentration rather than the free, unbound fraction that actually reaches the target, overstating the safety window. The same habit appears on a cardiotoxicity task (de_12, 0%). The pathways most associated with a three-level toxicity-concern score must be found with a rank test across all three levels, but models collapse the score to a yes/no label and rank pathways by a binary test, so their top pathways diverge from the correct set and they miss the potassium-channel mechanism the data supports.

Translational efficacy (49%) leans on reading microscopy images, and that reading, not the downstream judgment, is the bottleneck (perception is 62% of its scientific errors). Models do well when the evidence is a table, such as calling a genotype-by-treatment rescue (HAI2027_MM_FIG6C, 89%). When they must score an image they misjudge what they see, and because the image is the last readout before a go/no-go, the misread becomes the wrong decision. On a two-marker immunostain (HAI2027_MM_FIG2E, 42%) they read faint, near-vehicle single-agent staining as a real signal and credit each monotherapy with activity, endorsing a single-agent program when only the drug combination is active. The organoid viability readouts are the hardest tasks in the benchmark (as low as 6%). When apoptotic debris is counted as living organoids

(organoid_moa_cytotoxic), the mechanism call inverts, and a compound that is in fact killing the tumor cells is reported as

inert or growth-promoting and advanced on that false basis.

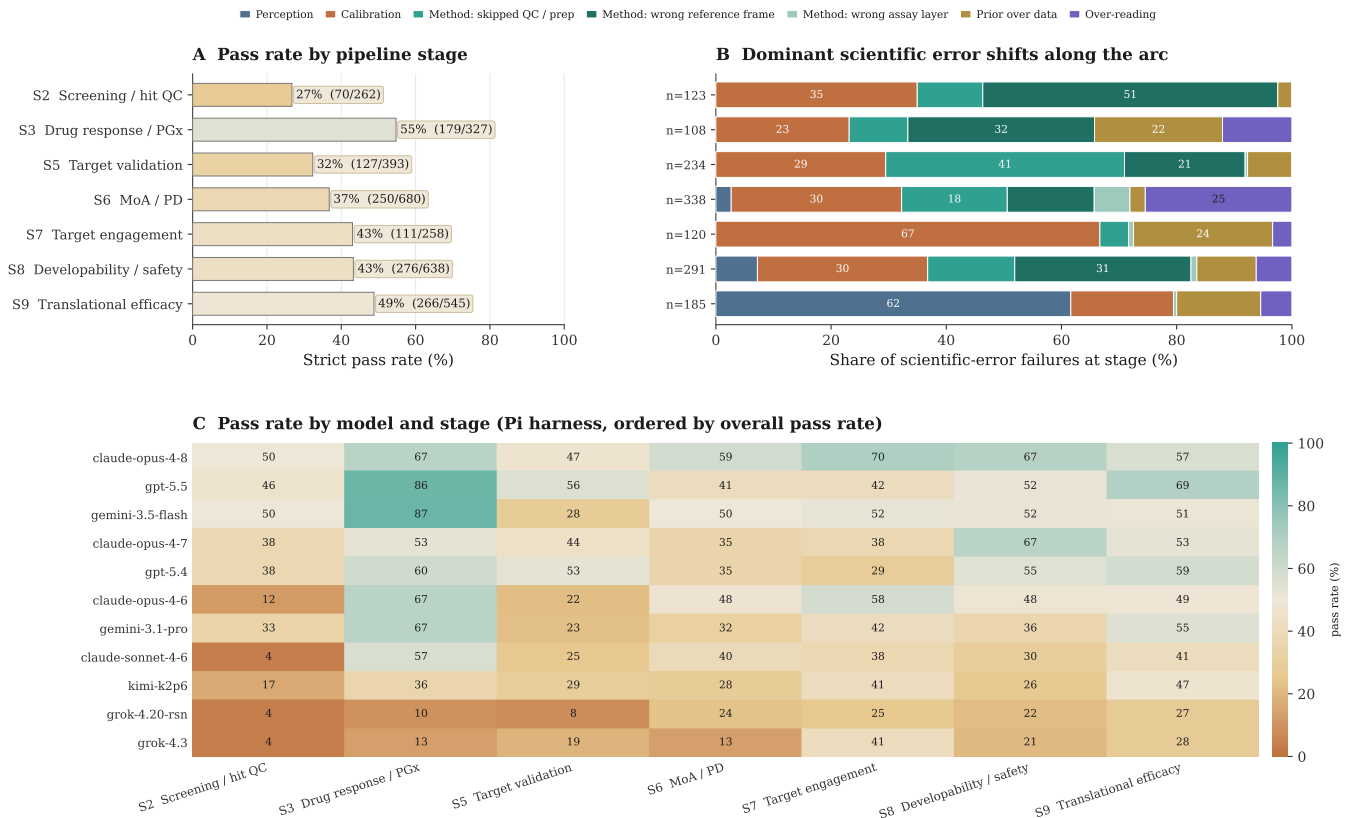


Figure 5: The dominant failure mode shifts across pipeline stages. Failing Pi-harness runs placed along the preclinical pharmacology arc. S1 disease-model and S4 target-genetics stages are omitted, since S1 is backed by too few evaluations and S4 is deferred to GenomicsBench. **(A)** Strict pass rate by stage. The two lowest are the early data-handling stages S2 (screening/QC) and S5 (causal target validation). **(B)** Scientific-error composition by stage, restricted to failures carrying a scientific error and using the same seven categories as Figure 4A. The dominant error shifts characteristically along the arc, from wrong reference frame at screening (S2) to skipped QC at target validation (S5), calibration at target engagement (S7), and perception at the image-heavy translational stage (S9), while calibration persists as a thread at every stage. **(C)** Strict pass rate for every model at each stage, ordered by overall pass rate. The overall ordering is largely preserved across stages, with two robust rearrangements: gemini-3.5-flash and gpt-5.5 lead at drug response (S3), and gemini-3.5-flash drops sharply at causal target validation (S5).

Some task categories remain far harder than others

Tasks become more difficult when they require nuanced biological reasoning and judgment. Grouping the 100 evaluations by task type and restricting to the 14 categories with at least four evaluations (86 of the 100 evaluations), the mean Pi-harness pass rate across the 11 models spans a 4.1-fold range, from 71% on exposure and PK tasks to 19% on hit prioritization and 17% on cheminformatics (Figure 6). Performance separates by the structure of the required answer rather than by subject matter. The four highest categories, exposure-PK (71%), biomarker discovery (56%), safety assessment (55%), and target engagement (52%), each resolve to a determinate target and average 58%. In each, the answer is either a num-

ber computed and compared against a threshold or a single entity identified by the evidence (a biomarker, or a binary safety or engagement call). The four lowest categories, target dependency (30%), differential response (28%), hit prioritization (19%), and cheminformatics (17%), instead require selecting a calibrated subset from many candidates with no unique answer, and average 24%. These tasks ask which genes are selective dependencies, whether a response is shared or lineage-selective, which screen hits to advance, and which substructures are genuine safety liabilities. The two groups differ by 34 percentage points (2.5-fold), and the ordering matches the failure taxonomy, in which tasks with a definite target are tractable for today’s models while those requiring a calibrated cut under uncertainty are not.




Figure 6: Some task categories remain far harder than others. Pi-harness pass rate by task type (categories with at least four evaluations), averaged across the 11 models. Difficulty tracks the kind of judgment required, from determinate quantitative tasks (exposure-PK) down to open-ended selection tasks (hit prioritization, cheminformatics).

Splitting performance by model and task type (Figure 7) shows that the two effects are largely separable. Within a model (a row of the matrix), overall model strength dominates, with the strongest models (claude-opus-4-8, gpt-5.5) scoring high across most categories and the weakest (grok-4.3,

grok-4.20-reasoning) scoring low across all of them. Within a task type (a column), difficulty imposes a ceiling that a stronger model does not overcome. On the two hardest categories, hit prioritization and cheminformatics, no model exceeds 45%, indicating a limit shared across current models.




Figure 7: Pass rate by model and task type. Strict Pi-harness pass rate for each of the 11 models (rows, ordered by overall pass rate) across the task categories with at least four evaluations (columns, ordered easiest to hardest). Model quality sets the level across a row, but the hardest categories on the right stay low for every model.

Among strong models, accuracy tracks process discipline rather than effort

Pass rate is only weakly and non-monotonically related to expended effort (Figure 8). Across the 11 models, mean agent steps and strict pass rate correlate at Spearman $\rho = 0.54$, dropping to $\rho = 0.32$ when the two low-engagement Grok models are excluded. grok-4.3 (mean \$0.02, median 6 steps) and grok-4.20-reasoning (\$0.11, median 11 steps) score lowest, at 18.3% and 19.7%, and under-investment drives 32% and 17% of their failures, against 0–1% for every other model (Figure 4C). Above the floor, effort and accuracy decouple. gemini-3.5-flash takes the most steps (median 43, $\approx 2.5\times$ opus-4-8's 17) but reaches only 51.3%, below opus-4-8's 59.3%. gpt-5.5 beats gpt-5.4 (55.3% vs. 49.7%) with fewer steps (median 15 vs. 24). claude-opus-4-6 reaches 44.7% at a median of 10 steps, above claude-sonnet-4-6 (36.0%) and kimi-k2p6 (29.7%) despite their higher step counts.

Strong models differ mainly in the composition of their failures, not in their analytic ceiling (Figure 4C). The top Claude models fail almost entirely through engaged-but-incorrect analysis. Analytical errors account for 73–77% of their failures, leaving 23–27% non-analytical (claude-opus-4-6 cleanest at 23.0% of $n = 165$ failures, opus-4-7 at 26.8%, opus-4-8 at 26.7%). The GPT models reach comparable accuracy by a more process-laden route. Weak commitment (gpt-5.5 12%, gpt-5.4 16%) and abandoned-correct answers (17% and 15%) raise their non-analytical share to 39.7% (gpt-5.5, $n = 126$) and 41.8% (gpt-5.4, $n = 146$). Their headroom therefore lies in

finalization, verifying the read and committing to the answer the analysis already supports, not in raw analytic capability. Under-investment is effectively absent from every strong model ($\leq 1\%$ of failures) and appears only in the two cheapest.

These quantitative patterns correspond to identifiable behaviors on individual tasks. The most accurate runs reason through the scientific intent, draw a calibrated cut, verify against the data, and commit to one answer. The pattern is clearest on hard tasks, where most runs fail but the strongest still succeed, and each hinges on a concrete drug-program decision. For instance, Bosutinib is an approved BCR-ABL/SRC inhibitor, but a program must credit its activity to the targets it actually engages in the chemoproteomics panel. On this task (32% of runs passing), claude-opus-4-8 overrode the clinical prior and committed to the non-canonical profile the panel supports, while other models anchored to the label returned the canonical ABL/SRC answer. The crizotinib control (21%) asks whether a separate broad-spectrum control compound justifies dropping genuinely engaged targets in a live-cell NanoBRET screen. It does not, and claude-opus-4-8 kept five engaged targets that other models discarded behind an invented quality-control gate, an error that would have removed real hits from medicinal-chemistry follow-up. The EGFR/phospho-MEK task (44%) asks whether EGFR productively drives downstream MEK signaling in a BRAF-mutant colorectal-cancer organoid, the call that decides whether the compound advances as on-pathway. gpt-5.5 quantified the co-positive cell share in code rather than calling it from total marker abundance, reaching the supported conclusion.




Figure 8: Spending more does not buy accuracy. Strict Pi-harness pass rate against mean cost (A) and mean agent steps (B) per evaluation (claude-opus-4-7 omitted, as its cost was not logged). The dashed line is the Pareto frontier. The cheapest, fewest-step models score lowest, but above a modest budget effort and accuracy decouple. The highest-step model (gemini-3.5-flash) is not the most accurate, and claude-opus-4-8 leads at moderate effort.

Failures translate directly into wrong go/no-go decisions

The evaluations that most resemble a real pipeline gate are the program decisions, which hand the agent several assay readouts at once and ask for a single advance, hold, or prioritize call. The current suite has seven such advancement decisions, and across 230 runs models pass only 35% of them (Figure 9). Models perform well when the evidence package is clean and consistent, but struggle on more challenging gates, either advancing weak candidates or discarding strong ones.

Clean correct. The organoid program go/no-go is the clean case (Figure 9). Given a primary screen, a mechanism image, a tox counterscreen, and DMPK exposure for eight blinded compounds, the supported action is to advance the one compound that shows the desired cytostatic mechanism together with an adequate safety margin and covering exposure, and most models do (82%).

Wrong synthesis. On the decision-gate task the agent must assemble a set of models with convergent support across several endpoints, but models rank on a single synergy metric and nominate a model that is an artifact of that one number (27%), so the program's next experiments are built on the wrong models.

False-go. On the safety advancement task the supported call is that no candidate moves forward, because none pairs activity with healthy-tissue sparing, yet models advance one on activ-

ity alone or on a relative tumour-versus-healthy index (36%), putting an unsafe compound into the pipeline.

False no-go. The reverse error also appears. On the long-term combination task, models read a strong colony-stain reduction as a complete kill and drop a model the combination still supports (HAI2027_FIG1_DUR01, 48%).

Leading the benchmark does not predict making the better program call. Six multi-evidence advancement decisions (18 runs each) yields an order almost unrelated to overall accuracy (Spearman $\rho = 0.08$; Figure 10A). claude-opus-4-8, first overall at 60%, ranks last on these decisions at 22% (4/18). It clears both single-compound calls, advancing the one qualifying organoid compound (2/3) and holding when no candidate is safe (2/3), but it fails all four decisions that instead ask which preclinical models (cell lines and PDX models) should anchor the next program stage (the single-agent, priority, long-term, and decision-gate sets in Figure 10B), passing none of them (0/3 each) because it over-ranks on a single metric or carries too many candidates forward. gpt-5.5 is the only model strong on both kinds of task (15/18, 83%), committing to a defensible prioritized set rather than a full ranking. kimi-k2p6, ninth overall at 31%, rises to second here (12/17, 71%): its careful per-model computation pays off when the task rewards naming the supported models, though it still fails the restraint call by advancing a candidate on the no-go safety task (0/3). The two Grok models stay at the bottom (5/18 each) by answering with little of the multi-evidence work the decision requires.

Program decisions integrate several assay readouts into one go/no-go

Across 7 such advancement evaluations (230 runs) models pass only 35%; the same call can fail by advancing too much or too little.

CORRECT	Advance one compound, or none Colorectal-cancer organoid program. Inputs: primary screen, mechanism image, tox counterscreen, DMPK exposure. Advance a compound only if the desired mechanism, an adequate safety margin, and covering exposure all hold.	DATA-SUPPORTED CALL Advance CMP-05, the one compound that satisfies all three criteria together.	MOSTLY CORRECT (82% PASS) The tempting shortcut, ranking by potency or by total exposure, advances the wrong compound.
	Pick the next decision-gate models NSCLC mTOR + WEE1 combination. The gate set must span models with confirmed multi-endpoint support, the strongest responder still lacking confirmation, and the strongest responder outside the primary genotype.	DATA-SUPPORTED CALL Build the gate on convergent multi-endpoint support (e.g. PDX239 for the unconfirmed slot).	COMMON FAILURE (27% PASS) Rank on a single synergy metric and nominate an artifact model (H2009), so the wrong models gate the program.
FALSE-GO	Advance only if safe and active Blinded intestinal multiplex-IF package. A candidate may advance only with paired evidence of activity and healthy-tissue sparing.	DATA-SUPPORTED CALL No candidate advances; the paired activity-and-safety evidence fails for every candidate.	COMMON FAILURE (36% PASS) Advance on activity alone or a relative tumour-vs-healthy index, pushing an unsafe candidate forward.

Figure 9: Program decisions turn analysis errors into wrong go/no-go calls. Three multi-evidence advancement evaluations, each giving the agent several assay readouts and asking for one advance, hold, or prioritize call. For each we show the evidence and the question, the data-supported call, and the common model error with its program consequence. The organoid go/no-go is passed by most models (82%); the decision-gate and safety-advancement tasks fail in opposite directions, advancing the wrong models (27%) or an unsafe candidate (36%).




Figure 10: Aggregate accuracy does not predict program-decision competence. Pass rates on the six multi-evidence advancement decisions (Pi harness, 18 runs per model). (A) Strict pass rate on all 100 evaluations versus on the six program decisions; the integrative go/no-go calls reorder the leaderboard almost completely (Spearman $\rho = 0.08$), with the overall-best model, claude-opus-4-8, falling to last and kimi-k2p6 rising from ninth to second. (B) Per-decision pass rate by model. The two single-compound calls (advance one, or hold; marked ●) are passed broadly, but the four tasks that require nominating the supported subset of models (marked with a diamond) separate the field; claude-opus-4-8 fails all four.

The agent harness has a surprising impact on model performance

Although a harness serves merely as a wrapper around a fixed model, utilizing the same checkpoint, task text, and answer-submission protocol. We found that it exerts a surprisingly large and consistent effect on overall performance. In within model matched comparisons, which provided the only unfounded contrast since each harness evaluated a distinct set of models, Pi outperformed Claude Code by 4.4 percentage points (95% CI: 1.2–7.8) on the shared Opus models. Similarly, Pi outperformed OpenAI Codex by 5.5 percentage points (95% CI: 1.2–9.8) on the shared GPT models. Notably, Pi achieved a higher score on every shared model, rather than merely winning on average.

The magnitude of this improvement is comparable to a model-generation upgrade. For instance, transitioning from Opus version 4.6 to 4.7 yielded only a 4.6-point increase on Pi. Consequently, the choice of harness is roughly as consequential

as upgrading to a newer model version. Furthermore, this advantage was relatively uniform across both image-based and non-image tasks, indicating that the performance gains were not primarily attributable to visual access. For context, Claude Code rendered figures to the model in 100% of image-bearing tasks compared to 84% for Pi, whereas OpenAI Codex, which lacks vision capabilities, never rendered figures.

Instead, the residual performance differences can be traced to subtler scaffolding properties. For example, Pi’s lean, three-tool loop failed on only 0.6% of runs—compared to 1.8% for Claude Code, and reached an answer in fewer steps. Conversely, OpenAI Codex’s lack of a vision pathway further penalized its performance on figure-decisive tasks. Although these effects are modest in absolute terms and present wide confidence intervals, they are directionally robust and persist across different model families. Ultimately, an implementation detail nominally orthogonal to the model itself can shift the endpoint pass rate by a margin equivalent to a full model generation.




Figure 11: Harness choice is as consequential as a model-generation upgrade. (A) Task-matched pass rate on Pi versus the alternate harness for each model pair. Points are matched on the shared evaluation set; the labeled gap is the within-model Pi advantage in percentage points. (B) Head-to-head comparison of the Pi advantage over an alternate harness (terracotta) against the gain from a model-generation upgrade on Pi (grey). A harness substitution shifts pass rate by a margin comparable to upgrading one model generation, so reported scores should be interpreted as a property of the model-plus-harness, not the model alone.

Discussion

TxBench-PP measures a narrow but important part of drug discovery: whether an agent can turn preclinical pharmacology evidence into a defensible local decision. This scope is intentionally smaller than end-to-end therapeutic discovery. It is a small-molecule task-group benchmark within the broader TherapeuticsBench roadmap, not a claim to cover every stage or modality. It does not yet test target genetics, full disease-model selection, primary screening operations, broad clinical translation, or non-small-molecule modalities. The narrower scope makes the benchmark verifiable while still covering decisions that directly affect compound advancement.

The central hypothesis is that current agents will often operate the tools of preclinical analysis but fail at the scientific judgment layer. If a model selects a mechanism from a drug label, trusts a contaminated target-engagement signal, ignores survivor bias, or ranks by total exposure rather than unbound coverage, it may produce a fluent and computationally elaborate answer that is not the pharmacology decision supported by the data.

[**TODO:** After results are available, add a paragraph comparing endpoint pass rates, partial-credit or field-level scores if available, and trajectory review. Emphasize what capability is missing and which tasks show meaningful progress.]

Future versions can extend the benchmark upstream to target rationale and primary screening, sideways to pharmacogenomics and disease-model characterization, and across modalities to antibodies, ADCs, degraders, oligonucleotides, cell therapies, and gene therapies. A complementary long-horizon TherapeuticsBench track can follow program stories across discovery, development, and translation. For the current paper, the clean claim should remain: TxBench-PP is a verifiable task-group benchmark for small-molecule preclinical pharmacology decisions.

Methods

Task construction

Candidate evaluations are derived from realistic preclinical pharmacology workflow states. A task is retained when the target conclusion can be recovered from the supplied files, expressed through a constrained final answer, and graded deterministically. Tasks should be revised or removed when the prompt over-specifies the method, when multiple reasonable

analyses produce incompatible decisions, or when the grader cannot distinguish the supported pharmacology conclusion from a plausible but unsupported answer.

Agent runs and aggregation

[**TODO:** Insert model-harness list, number of attempts per evaluation, execution environment, parsing rules, and how invalid or missing outputs are counted.]

Statistical analysis

[**TODO:** Insert primary metric and confidence interval convention. If following SpatialBench/EpiBench, report evaluation-weighted endpoint pass rate, raw pass counts for interpretability, and Student t confidence intervals over evaluation-level replicate means.]

Data availability

[**TODO:** Insert public repository, public example evaluations, result summaries, and trajectory availability language.]

References

- [1] Zecha, J. et al. Decrypting drug actions and kinase functions by dose-resolved phosphoproteomics. *Science* (2023).
- [2] Eckert, M. A. et al. decryptE: dose-resolved expression proteomics for compound mechanism analysis. *Manuscript / dataset reference to verify before submission* (2025).
- [3] Reinecke, M. et al. Kinobeads profiling for chemoproteomic target engagement. *Manuscript / dataset reference to verify before submission* (2024).
- [4] Igarashi, Y. et al. Open TG-GATEs: a large-scale toxicogenomics database. *Nucleic Acids Research* 43, D921–D927 (2015).
- [5] Ganter, B. et al. DrugMatrix toxicogenomics resource. *Pharmacogenomics* 6, 715–727 (2005).
- [6] Corsello, S. M. et al. Discovering the anticancer potential of non-oncology drugs by systematic viability profiling. *Nature Cancer* 1, 235–248 (2020).
- [7] Subramanian, A. et al. A next generation connectivity map: L1000 platform and the first 1,000,000 profiles. *Cell* 171, 1437–1452 (2017).
- [8] Chandrasekaran, S. N. et al. JUMP Cell Painting dataset. *Manuscript / resource reference to verify before submission* (2024).