

TxBench-PP: Verifiable Evaluation of AI Agents on Preclinical Pharmacology

Hannah Le^{1,*}, Ramesh Ramasamy^{1,*}, Alex Urrutia^{1,*}, Mahsa Yazdani^{1,*}, Tim Proctor¹, Kenny Workman¹

¹LatchBio, San Francisco, CA, USA

*These authors contributed equally.

Correspondence: kenny@latch.bio

ABSTRACT

We introduce TxBench-PP, a verifiable benchmark for evaluating AI agents on small-molecule preclinical pharmacology. TxBench-PP is the first focused slice of a broader TherapeuticsBench effort organized across drug-discovery stages and therapeutic modalities. It targets recurring cluster-style decisions: whether agents can recover decision-relevant conclusions from supplied assay evidence rather than recalling drug labels, ranking by a single metric, or applying generic checklists. The current task inventory maps 70 evaluations to a nine-stage Assays Roadmap. Seven stages have coverage, with strongest current coverage in causal target validation and perturbational mechanism-of-action / pharmacodynamic reasoning, single-source coverage in screening, pharmacogenomics, safety, and translational efficacy, and a near-empty compound-target characterization stage. Agents receive realistic workflow snapshots, inspect files in a coding environment, and return structured answers graded deterministically. [TODO: Replace this final paragraph with model-run results: number of evaluations, model-harness pairs, trajectories, top pass rates, and the main failure pattern.]

TxBench-PP in TherapeuticsBench

TxBench-PP is the current small-molecule preclinical pharmacology slice of a broader TherapeuticsBench roadmap. The full benchmark family is organized by drug-discovery stage and therapeutic modality; this paper focuses on repeated, verifiable cluster decisions in small-molecule programs. These decisions sit between initial discovery and clinical development: whether a molecule engages the intended biology, whether the evidence supports its mechanism, whether exposure can cover the relevant potency range, and whether safety signals block advancement.

TxBench-PP is a focused small-molecule slice of TherapeuticsBench

The broader roadmap is organized by stage and modality; the current benchmark tests recurring cluster decisions.




Figure 1: TxBench-PP as a TherapeuticsBench slice. TherapeuticsBench is intended to span discovery, development, and translation across small molecules, biologics, and oligonucleotides. TxBench-PP covers the small-molecule preclinical pharmacology slice as a cluster benchmark: repeated local decisions across assay families and compounds. Program-story benchmarks that follow one therapeutic program across discovery, development, and translation are reserved for future work.

Introduction

In preclinical drug discovery, experiments are slow and expensive, but many program-limiting decisions are interpretive: which signal is real, which compound should advance, and which liability should block a program. Preclinical pharmacology converts assay evidence into these decisions. A compound can have an attractive potency value and still be a poor candidate because the apparent mechanism is an artifact, the engaged protein is a complex contaminant, the safety signal is protocol-specific, or total exposure does not translate into unbound target coverage.

AI agents can inspect files, write analysis scripts, and invoke scientific software, but their conclusions often remain vulnerable to plausible shortcuts. In drug discovery, those shortcuts have familiar forms: importing the known mechanism of a blinded compound, trusting the largest apparent effect without quality control, ranking by potency or total exposure alone, or applying a textbook safety threshold outside the protocol where it was defined.

We introduce TxBench-PP, a benchmark for small-molecule preclinical pharmacology decisions. It should be read as a focused cluster benchmark within the broader TherapeuticsBench roadmap rather than as a complete test of drug discovery. Each evaluation supplies the agent with realistic analysis artifacts from one stage of the preclinical pharmacology evidence chain and asks for a structured, deterministically gradable conclusion. The benchmark tests whether an agent can recover the result supported by the provided evidence, not whether it can recall a drug class, run a default workflow, or optimize for a single number.

The current benchmark maps 70 evaluations to a nine-stage Assays Roadmap. Coverage is strongest in causal target validation and perturbational mechanism-of-action / pharmacodynamic reasoning; the clearest build gap is compound-target characterization. The underlying sources still matter: several covered stages are dominated by a small number of datasets, so diversifying assay sources is a central part of the benchmark roadmap. **[TODO: Add final run summary once result aggregation is complete.]**

Benchmark Construction

To construct TxBench-PP, we decompose preclinical pharmacology workflows into short, gradeable scientific decisions. Examples include identifying a pathway-coherent mechanism from dose-resolved phosphoproteomics, separating true chemoproteomic target engagement from complex contaminants, reconciling conflicting clearance measurements, and integrating blood chemistry, pathology, and survival into an in vivo safety call.

This makes TxBench-PP a cluster-style benchmark: the same kind of decision is tested repeatedly across assay contexts, compounds, and evidence sources. Future TherapeuticsBench story evaluations can test the complementary skill of following one therapeutic program across discovery, development, and translation. TxBench-PP instead asks whether agents can make the local scientific decisions that those longer program stories depend on.

Each task includes a snapshot of assay outputs immediately before a target decision, metadata needed to interpret the files, a high-level task description, and a deterministic grader. Tasks specify what result should be recovered rather than prescribing the exact analysis method. Procedural details are included only when they are part of the scientific decision, such as whether free or total exposure should be compared to a pharmacologic threshold.

The benchmark follows three design criteria. First, tasks must be verifiable: answers are checked by structured label matching, numerical tolerances, rank-order checks, or all-of field comparisons. Second, tasks should be decision-relevant: the answer should correspond to a real go/no-go or prioritization judgment in preclinical pharmacology. Third, tasks should resist shortcuts: agents should need to inspect the supplied evidence, account for quality issues, and avoid relying on memorized drug labels or generic checklists.

We separately track decision-blocking gaps: issues that should change the program decision if recognized, such as a tox-species mismatch, insufficient free-drug coverage, a hook effect, a contaminated chemoproteomic hit, or survivor bias in pathology sampling. These labels are intended to support diagnostic analysis after endpoint grading.

Evaluation Inventory

The current suite is organized around the Assays Roadmap rather than around source datasets. Each evaluation is assigned to the roadmap stage that best matches its primary biological skill. The current inventory includes 70 evaluations across seven of the nine stages. S1 disease/model characterization and S4 human genetic target support are deferred for this release; S4 is better handled by GenomicsBench.

The roadmap view changes the benchmark story. Coverage is broad by stage count, but narrow by dataset diversity. Causal target validation and perturbational MoA/PD are ready for v1. Screening, pharmacogenomics, safety, and translational efficacy have useful coverage but are single-source dominated. The largest in-scope gap is S7 compound-target characterization: binding, potency, selectivity, and cellular target engagement are almost absent.




Figure 2: Evaluation inventory by Assays Roadmap stage. Current TxBench-PP evaluations are mapped to the roadmap stage that best captures the biological skill being tested. The inventory is broad across stages but uneven across assay families: S5 and S6 are ready for v1, S7 is the clearest build gap, and S2, S3, S8, and S9 need source diversification.

Table 1: Roadmap-stage inventory. Counts are from the Assays Roadmap coverage mapping of 70 current evaluations. “Ready” stages are sufficient for v1; “Diversify” stages have useful but source-narrow coverage; “Build” marks the largest missing pillar.

Roadmap stage	Evals	Status	Assay emphasis and next gap
S1. Disease / model characterization	0	Deferred	Model subtype, expression, mutation/CNV context, pathway state, IHC/IF, RNAscope, and model QC are out of the current build queue.
S2. Primary screening / hit confirmation	7	Diversify	PRISM-heavy. Add biochemical or reporter HTS, Z-prime and signal-window plate QC, counterscreens, and frequent-hitter filtering.
S3. Drug response / pharmacogenomics	11	Diversify	Mostly PRISM. Add GDSC, CTRP, CCLE, Beat AML, organoid screens, single-cell perturbation, and censored-curve handling.
S4. Human genetic target support	0	Deferred	GWAS, QTL, colocalization, Mendelian randomization, and tissue-of-action matching are assigned to GenomicsBench.
S5. Causal target validation	12	Ready	Strong v1 coverage. Optional expansion: Perturb-CITE-seq, rescue assays, variant-function classification, and FACS sort-bin screens.
S6. Perturbational MoA / PD	17	Ready	Strong v1 coverage. Optional expansion: LINCSC/CMAP, DRUG-seq, Seahorse, cytokine panels, reporter dose-response, and washout PD.
S7. Compound-target characterization	1	Build	Highest-priority gap. Add ChEMBL/BindingDB potency curation, selectivity panels, CETSA/NanoBRET, Kinobeads, and biochemical-to-cellular potency comparison.
S8. Developability / safety	16	Diversify	Cardiotox dominated. Add measured hERG, DILI, Ames/micronucleus, permeability, metabolic stability, CYP/DDI, Tox21/ToxCast, and therapeutic-index reasoning.
S9. Translational efficacy	6	Diversify	Single NSCLC source. Add organoid/patient concordance, PDX pharmacology, immune co-culture, ctDNA, responder-fraction, and exposure-response PK/PD.

Cross-stage gap. The roadmap also identifies a missing class of end-to-end go/no-go evaluations: tasks that require an agent to weigh efficacy, selectivity, target engagement, exposure, and safety together. The current inventory covers many of these component decisions, but future TxBench-PP tasks should test multi-stage synthesis across at least three evidence blocks.

Results

No agent reliably recovers preclinical pharmacology decisions

[**TODO:** Topline paragraph from the model leaderboard (fig1_leaderboard): state the scope (11 model families, 100 evaluations $\times$ 3 rounds on the Pi harness), the pass-rate definition, and the ordering from the leader (claude-opus-4-8, $\sim$ 60%) down to the floor (grok-4.3, $\sim$ 19%), with raw counts and confidence intervals. Emphasize that even the best system leaves roughly 40% of decisions unrecovered.]

Most failures map to a compact failure-mode taxonomy

To understand why agents fail, we read every failing (model, evaluation) run on the Pi harness and labeled it against the evaluation’s documented scientific intent, resulting in 1,834 labeled failures. Each failure is classified along two complementary axes. The *scientific error* captures what is wrong with the answer, and the *behavioral cause* captures how the model arrived there (Figure 3).

The scientific axis comprises five recurring error families (Figure 3A, restricted to the $n = 1,436$ failures that contain a scientific error). The largest is method, which resolves into two main sub-modes plus a small tail. A *skipped-QC* method error (17%) omits a required preprocessing or quality-control step. In one case, a model aggregated CRISPR guides to gene-level scores without first dropping guides that map to multiple loci, so paralog-family artifacts flooded the essential-gene list. A *wrong-reference-frame* method error (21%) runs the pipeline but compares against the wrong baseline, comparator, or normalization. A small *wrong-assay-layer* tail (2%) reads the wrong assay feature or modality entirely. A *calibration* error (31%) draws a cutoff a domain expert would set differently. In one competition-binding panel, models carried forward all 33 nonzero kinases when only the 8 strong binders above a clear affinity gap were supported. A *perception* error (10%) misreads a raw image, for example counting apoptotic debris as viable organoids or mis-segmenting two-channel colocalization. *Prior over data* (10%) answers from textbook or famous biology against the staged evidence. One model computed the strongest data-supported biomarker, a vanadium compound correlated with SLC26A2 at $r = -0.48$, and then discarded

it for the textbook elesclomol-FDX1 cuproptosis pair. *Over-reading* (10%) states a stronger claim than the data supports, for example reporting tumor regression from a sub-baseline median when the per-animal measurements show only growth arrest. Method and calibration together account for 71% of scientific errors.

The behavioral axis records how the failure happened, over all $n = 1,834$ runs (Figure 3B). Two-thirds (67%) are *engaged, analytically wrong*. The model did the work in good faith and made an analytical error. The remaining third splits into one evidence-handling failure and three process failures. *Prior-driven override* (6%) is the one behavior in which the model engaged the data, often computed the supporting result, and then deferred to a textbook or famous-biology prior. It is the behavioral signature of the scientific *prior over data* error, and the two coincide almost exactly. *Weak commitment* (9%) performs the right analysis but submits the wrong final answer. In one run a model’s reasoning correctly identified CTNNB1 as the supported dependency while its output named E2F3. *Abandoned-correct* (12%) is the sharpest case. The model demonstrably had a passing result within reach and then committed to a worse one. *Under-investment* (6%) reaches an answer with little tool use or verification, for example a mitotic-kinase family call written from prior knowledge after merely listing the compounds. Per-run effort separates these modes. Under-investment is the only one written with markedly less work (median 6 reasoning steps against 14–17 for the others). Abandoned-correct and prior-driven override occur in the *most* effortful runs (median 16 steps), so they are not relieved by additional compute. That a third of failures are non-analytical, and that abandoned-correct alone covers 212 runs in which the right result was within reach, implies that a meaningful fraction is recoverable through better verification and final-answer selection rather than added capability.

The failures have a unifying frame. These are fluency-optimized systems applied to a domain that rewards adversarial restraint. The fluent move is confident, coherent, and consistent with prior knowledge, while the right experimental move defers to the data over the textbook, draws conservative cuts, and distrusts its own first answer. TxBench-PP’s traps sit where these diverge, so the failures are a systematic readout of that misalignment. The models have the knowledge and the tooling, and often compute the right intermediate result, but they lack the epistemic stance of an experimentalist, namely prior-skepticism, calibrated commitment, and self-verification.




Figure 3: A two-axis failure taxonomy. All 1,834 failing (model, evaluation) runs on the Pi harness, labeled against each evaluation’s scientific intent and classified along two complementary axes. **(A)** The scientific-error axis (what is wrong with the answer) is dominated by method and calibration errors (71% combined), which are general analytical skills rather than domain knowledge; the method family resolves into a skipped-QC/preprocessing sub-mode and a larger wrong-reference-frame sub-mode. **(B)** The behavioral-cause axis (how the model failed) shows that a third of failures are non-analytical: a prior-driven override of the model’s own evidence, plus three process breakdowns (weak commitment, abandoned-correct, under-investment). **(C)** The behavioral mix is a model signature: under-investment dominates only the weakest model, while the strongest models fail almost entirely through engaged, analytically wrong reasoning. A scientific-error $\times$ behavioral-cause cross-tabulation is provided in the supplement (fig_fm_crosstab).

The dominant failure mode shifts characteristically across pipeline stages

Pass rates range from 27% at the hardest stage to 55% at the most forgiving (Figure 4). Walking the arc in order shows that each stage has characteristic strengths and weaknesses. We look at a representative easy and hard evaluation for each stage, with the fraction of passing runs in parentheses.

Primary screening and hit prioritization (27%) is the hardest stage. Models are competent at the mechanical steps, such as robustly collapsing noisy replicates (PRISM05, 58%). They fail by choosing the wrong reference frame for QC. On outlier-pool QC (PRISM01, 18%) they flag whole well-by-pool groups by an aggregate median residual and exclude most of the cytotoxic positive-control wells, discarding real biology as a technical artifact. On an oncology drugs repurposing prioritization task (PRISM_S2_11, 0%) they rank by breadth of killing and nominate a pan-cytotoxic compound rather than a selective hit.

Drug response and pharmacogenomics (55%) is one of the easier stages but exposes a consistent weakness in unsupervised mechanism recovery. Models do well on direct evidence checks, such as separating a causal mediator from a correlated biomarker (PRISM_S4_06, 76%). They fail when asked to cluster compounds by mechanism from their response profiles. Across the embedding tasks they use a potency-dominated representation, a linear PCA or a Euclidean metric on raw profiles, so the leading component encodes potency and the intended drug class never separates from the rest. The same error recurs across statin collapse (PRISM_UMAP_01, 12%), oncology MoA recovery (PRISM_UMAP_04, 48%), and glucocorticoid substructure (PRISM_UMAP_03, 69%), so it is a property of the stage rather than a single task.

Causal target validation (32%) is the second-hardest stage and the one that most rewards data hygiene and statistical rigor. Models recover genuine dependencies from a clean screen (HIT_05 replicate discordance, 69%). They routinely skip the library hygiene that CRISPR screens require. On multi-target guide contamination (GG_01, 16%) they aggregate guides to gene level without dropping multi-locus sgRNAs, so paralog-family false positives flood the dependency list. On the selective-set call (SEL_01, 6%) they ship thousands of genes behind a count-only filter with no effect-size floor.

Mechanism of action and PD (37%) is the largest stage, with 21 evaluations. Models reliably recover known program structure, such as assigning cells to the correct cell-cycle phase (clust_04, 84%). Their recurring failure is over-reaching on the mechanism, treating marginal or partial evidence as a strong, broad effect. On the anti-infective task (ANTIINF_01, 0%) they read near-noise pathway scores ($p \approx 0.05$) as

genuinely engaged mechanisms. On the mTOR fingerprint (TXBX_DM_v15E, 6%) they correctly call mTOR the primary target but then count phosphosites from too broad a substrate panel, sweeping in adjacent AKT-axis sites instead of restricting to the canonical mTOR set, which inflates the affected-site count to almost twice what the data supports (59 versus 31).

Compound-target engagement (43%) asks whether a compound truly hits its target and how selectively, and the bottleneck is calibration, not detection (67% of the stage's scientific errors). Models detect the binding signal reliably, for example confirming a clean direct target by thermal shift (thermal_shift_shared_direct_target, 91%). What they lack is a principled cutoff for how much engagement counts, so they err in both directions. On a kinase-binding panel (TXBX_KB_v11_DA1, 0%) they accept every weak binder as a real hit, and on a control task (CTRL01, 21%) they invent a stringency cutoff the data never justified and discard five genuinely engaged targets.

Developability and safety (43%) depends on judging a signal against the right reference, and that is where models slip. They make a confident call when the standard is unambiguous, such as the hepatotoxicity verdict (TXBX_TG_v11_A, 91%). On the hERG safety-margin task (de_14) they size the margin against total drug concentration rather than the free, unbound fraction that actually reaches the target, overstating the safety window. The same habit appears on a cardiotoxicity task (de_12, 0%). The pathways most associated with a three-level toxicity-concern score must be found with a rank test across all three levels, but models collapse the score to a yes/no label and rank pathways by a binary test, so their top pathways diverge from the correct set and they miss the potassium-channel mechanism the data supports.

Translational efficacy (49%) leans on reading microscopy images, and that reading, not the downstream judgment, is the bottleneck (perception is 62% of its scientific errors). Models do well when the evidence is a table, such as calling a genotype-by-treatment rescue (HAI2027_MM_FIG6C, 89%). When they must score an image they misjudge what they see, and because the image is the last readout before a go/no-go, the misread becomes the wrong decision. On a two-marker immunostain (HAI2027_MM_FIG2E, 42%) they read faint, near-vehicle single-agent staining as a real signal and credit each monotherapy with activity, endorsing a single-agent program when only the drug combination is active. The organoid viability readouts are the hardest tasks in the benchmark (as low as 6%). When apoptotic debris is counted as living organoids (organoid_moa_cytotoxic), the mechanism call inverts, and a compound that is in fact killing the tumor cells is reported as inert or growth-promoting and advanced on that false basis.




Figure 4: The dominant failure mode shifts across pipeline stages. Failing Pi-harness runs placed along the preclinical pharmacology arc. S1 disease-model and S4 target-genetics stages are omitted, since S1 is backed by too few evaluations and S4 is deferred to GenomicsBench. **(A)** Strict pass rate by stage. The two lowest are the early data-handling stages S2 (screening/QC) and S5 (causal target validation). **(B)** Scientific-error composition by stage, restricted to failures carrying a scientific error and using the same seven categories as Figure 3A. The dominant error shifts characteristically along the arc, from wrong reference frame at screening (S2) to skipped QC at target validation (S5), calibration at target engagement (S7), and perception at the image-heavy translational stage (S9), while calibration persists as a thread at every stage.

Some task categories remain far harder than others

Difficulty depends more on the kind of judgment a task requires than on its biological domain. Grouped by task type rather than by pipeline stage, mean pass rate across the 11 models spans roughly fourfold over the categories with at least four evaluations, from 71% on exposure and PK tasks to 19% on hit prioritization and 17% on cheminformatics (Figure 5). The high-scoring categories have a clear, well-defined answer, either a number computed and compared against a threshold (exposure and PK margins) or a single entity the evidence

points to, such as a biomarker or a binary safety or engagement call. The low-scoring categories instead ask the model to pick a calibrated subset from many candidates, with no single answer to recover. Examples are which screen hits to advance (hit prioritization), which genes are selective dependencies (target dependency), whether a response is shared or lineage-selective (differential response), and which substructures are genuine safety liabilities (cheminformatics). The pattern matches the failure taxonomy. Tasks with a definite target are tractable, while those that require committing to a calibrated cut under uncertainty are not.




Figure 5: Some task categories remain far harder than others. Pi-harness pass rate by task type (categories with at least four evaluations), averaged across the 11 models. Difficulty tracks the kind of judgment required, from determinate quantitative tasks (exposure-PK) down to open-ended selection tasks (hit prioritization, cheminformatics).

Splitting performance by model and task type (Figure 6) shows the two effects act largely independently. Across a row, overall model strength dominates, with the strongest models (claude-opus-4-8, gpt-5.5) high across most categories and the weakest (grok-4.3, grok-4.20-reasoning) low throughout.

Down a column, the task sets a ceiling that a better model does not overcome. On the two hardest categories, hit prioritization and cheminformatics, no model exceeds 45%, which points to a limit shared across current models,




Figure 6: Pass rate by model and task type. Strict Pi-harness pass rate for each of the 11 models (rows, ordered by overall pass rate) across the task categories with at least four evaluations (columns, ordered easiest to hardest). Model quality sets the level across a row, but the hardest categories on the right stay low for every model.

The strongest models win through disciplined scientific process

Pass rate is only weakly related to expended effort, and not monotonically. Across the 11 models, neither mean cost nor mean number of agent steps predicts accuracy above a low threshold (Figure 7). The two cheapest and fastest models, grok-4.3 (\$0.02 and 6 steps per evaluation) and grok-4.20-reasoning (\$0.11 and 11 steps), score lowest (19% and 20%), with failures dominated by under-investment (Figure 3C). Above that level the two quantities diverge. gemini-3.5-flash uses the most steps (median 42, roughly 2.5 times opus-4-8) but scores 52% against opus-4-8's 60%, gpt-5.5 exceeds gpt-5.4 with fewer steps (15 versus 24), and claude-opus-4-6 reaches 45% in about 10 steps, above claude-sonnet-4-6 and kimi-k2p6 at higher step counts.

The differences among the stronger models track process discipline, which is comparable across systems. The most accurate runs reason through the question's scientific intent, draw a calibrated cut, verify intermediate results against the data, and

commit to a single answer. claude-opus-4-8, the top model, defers to the staged evidence over priors. On the bosutinib task it set aside the compound's documented BCR-ABL and SRC activity and reported the non-canonical off-target engagement the data show, the distinction that determines whether a program attributes the compound's activity to its reputed targets or to the ones it actually engages. On the crizotinib control it applied only the threshold task specified, rather than the additional gate other models introduced that removes five genuine engagement hits and understates the compound's cellular target profile. gpt-5.5, the second-ranked model, is the best-calibrated on precision-sensitive tasks and verifies its reads before answering. On an EGFR and phospho-MEK colocalization task it established the panel layout and marker assignments before quantifying co-activation in code, the readout that decides whether the compound is advanced as on-pathway. Consistent with this, the strongest models fail predominantly through engaged but incorrect analysis rather than under-investment or weak commitment, which are concentrated in the lower-ranked models (Figure 3C).




Figure 7: Spending more does not buy accuracy. Strict Pi-harness pass rate against mean cost (A) and mean agent steps (B) per evaluation (claude-opus-4-7 omitted, as its cost was not logged). The dashed line is the Pareto frontier. The cheapest, fewest-step models score lowest, but above a modest budget effort and accuracy decouple. The highest-step model (gemini-3.5-flash) is not the most accurate, and claude-opus-4-8 leads at moderate effort.

Failures translate directly into wrong go/no-go decisions

[TODO: Use the downstream-consequence lists in _stage_failure_maps.md to tie representative failures to

concrete program errors (e.g. advancing a non-selective compound as selective, a wrong cytostatic-vs-cytotoxic call, a missed safety liability, a contaminated target-engagement hit carried forward). Ground each in an exact eval name once the public manifest is frozen.]

Representative task logic spans the preclinical pharmacology spine




Figure 8: Representative task logic. Across mechanism, target engagement, safety, and exposure decisions, the correct answer is the one supported by the supplied assay evidence, not the familiar shortcut.

The agent harness has a surprising impact on model performance

[**TODO:** From the harness-effect comparison (fig4_harness_effect): running the same model under different scaffolds (Pi vs Codex/CC) shifts the task-matched

pass rate by roughly 3–12 points, with some models far more harness-sensitive than others (e.g. claude-opus-4-7) even though the overall ordering is largely preserved. Note the implication for fair model comparison: reported scores are a property of the model-plus-harness, not the model alone.]

Discussion

TxBench-PP measures a narrow but important part of drug discovery: whether an agent can turn preclinical pharmacology evidence into a defensible local decision. This scope is intentionally smaller than end-to-end therapeutic discovery. It is a small-molecule cluster benchmark within the broader TherapeuticsBench roadmap, not a claim to cover every stage or modality. It does not yet test target genetics, full disease-model selection, primary screening operations, broad clinical translation, or non-small-molecule modalities. The narrower scope makes the benchmark verifiable while still covering decisions that directly affect compound advancement.

The central hypothesis is that current agents will often operate the tools of preclinical analysis but fail at the scientific judgment layer. If a model selects a mechanism from a drug label, trusts a contaminated target-engagement signal, ignores survivor bias, or ranks by total exposure rather than unbound coverage, it may produce a fluent and computationally elaborate answer that is not the pharmacology decision supported by the data.

[**TODO:** After results are available, add a paragraph comparing endpoint pass rates, partial-credit or field-level scores if available, and trajectory review. Emphasize what capability is missing and which tasks show meaningful progress.]

Future versions can extend the benchmark upstream to target rationale and primary screening, sideways to pharmacogenomics and disease-model characterization, and across modalities to antibodies, ADCs, degraders, oligonucleotides, cell therapies, and gene therapies. A complementary long-horizon TherapeuticsBench track can follow program stories across discovery, development, and translation. For the current paper, the clean claim should remain: TxBench-PP is a verifiable cluster benchmark for small-molecule preclinical pharmacology decisions.

Methods

Task construction

Candidate evaluations are derived from realistic preclinical pharmacology workflow states. A task is retained when the target conclusion can be recovered from the supplied files, expressed through a constrained final answer, and graded deterministically. Tasks should be revised or removed when the prompt over-specifies the method, when multiple reasonable

analyses produce incompatible decisions, or when the grader cannot distinguish the supported pharmacology conclusion from a plausible but unsupported answer.

Agent runs and aggregation

[**TODO:** Insert model-harness list, number of attempts per evaluation, execution environment, parsing rules, and how invalid or missing outputs are counted.]

Statistical analysis

[**TODO:** Insert primary metric and confidence interval convention. If following SpatialBench/EpiBench, report evaluation-weighted endpoint pass rate, raw pass counts for interpretability, and Student *t* confidence intervals over evaluation-level replicate means.]

Data availability

[**TODO:** Insert public repository, public example evaluations, result summaries, and trajectory availability language.]

References

- [1] Zecha, J. et al. Decrypting drug actions and kinase functions by dose-resolved phosphoproteomics. *Science* (2023).
- [2] Eckert, M. A. et al. decryptE: dose-resolved expression proteomics for compound mechanism analysis. *Manuscript / dataset reference to verify before submission* (2025).
- [3] Reinecke, M. et al. Kinobeads profiling for chemoproteomic target engagement. *Manuscript / dataset reference to verify before submission* (2024).
- [4] Igarashi, Y. et al. Open TG-GATEs: a large-scale toxicogenomics database. *Nucleic Acids Research* 43, D921–D927 (2015).
- [5] Ganter, B. et al. DrugMatrix toxicogenomics resource. *Pharmacogenomics* 6, 715–727 (2005).
- [6] Corsello, S. M. et al. Discovering the anticancer potential of non-oncology drugs by systematic viability profiling. *Nature Cancer* 1, 235–248 (2020).
- [7] Subramanian, A. et al. A next generation connectivity map: L1000 platform and the first 1,000,000 profiles. *Cell* 171, 1437–1452 (2017).
- [8] Chandrasekaran, S. N. et al. JUMP Cell Painting dataset. *Manuscript / resource reference to verify before submission* (2024).